

Stroke Prediction

Using machine learning to predict
the onset of stroke

Josh Wang
11/21/2022



Background



A stroke occurs when the blood flow to the brain is blocked. It is the **5th cause of death and a leading cause of disability** in the United States.



A stroke is an emergency situation, and it is important to be able to predict the onset of stroke. (Healthcare, insurance, etc.)



The Data (from Kaggle)

features



age



gender



**blood
glucose**



**smoking
status**



**heart
disease**

etc.



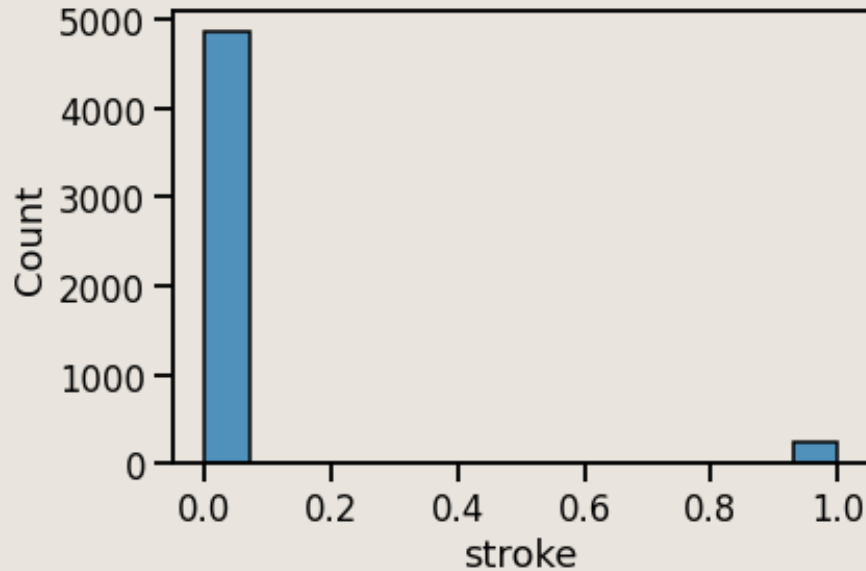
target



**stroke
(Y/N)**



The Data – Class Imbalance

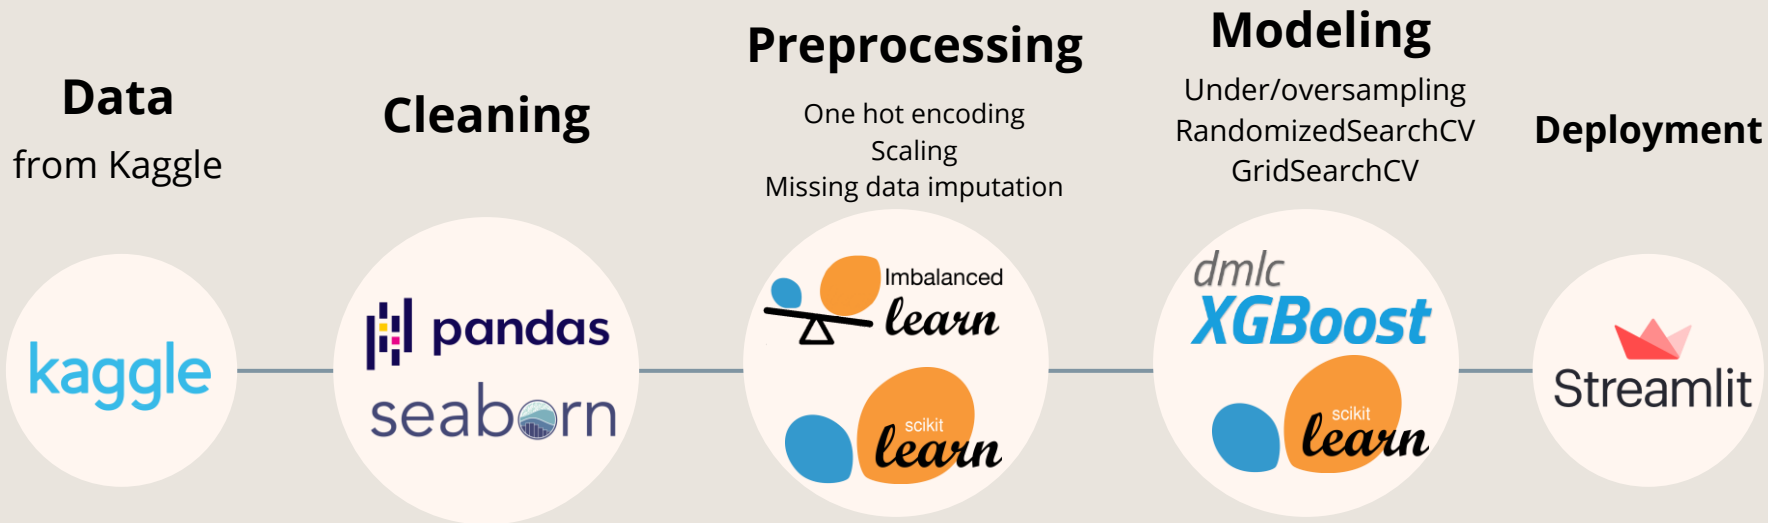


Data is heavily skewed.

Only 5% of the datapoints are positive for stroke.

{0: 95%, 1: 5%}

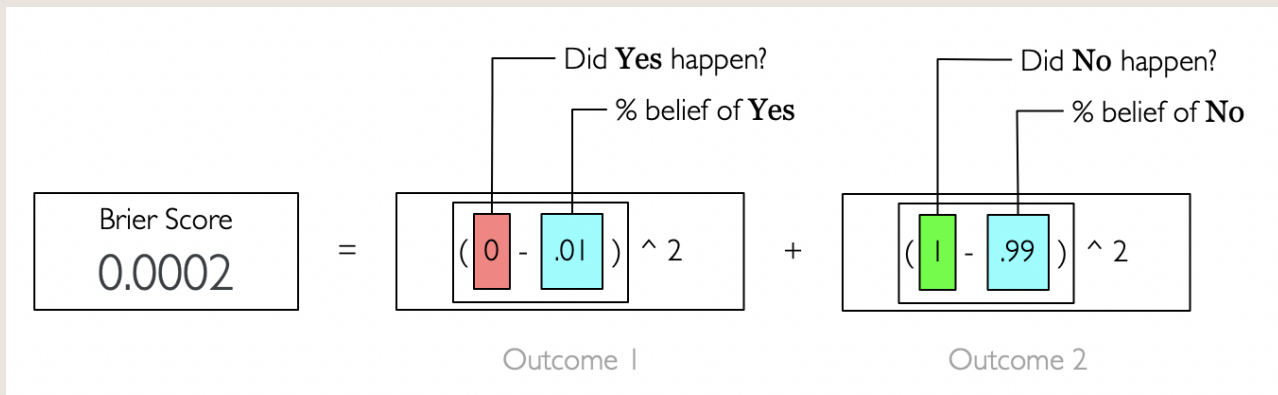
Methodology



Evaluation Metric – Brier Skill Score

Brier Score is the sum of the squares of the residuals. (similar to Mean Squared Error)

Measures the accuracy of probabilistic predictions (lower is better):



Brier Skill Score (BSS) = $1 - \text{BS}/\text{BS}^{\text{ref}}$ (higher is better)

BSS>0, model better than baseline

BSS=0, model has no skill

BSS<0, model worse than baseline

Results

Cross-validation scores (higher is better):


Model	Base Model BSS	Tuned Model BSS
Logistic Regression	0.075	0.076
kNN	0.013	0.049
RandomForest	0.034	0.077
XGBoost	0.021	0.075
SVC	0.007	0.083

Probabilistic outputs were calibrated with Platt Scaling when necessary.

SVC wins. The BSS of SVC on the holdout data is 0.113

Results

Feature coefficients from Logistic Regression:





Product Deployment: A Web App to Predict Stroke Onset

(powered by Streamlit)

(Click [here](#))



Image from Unsplash

Upload your input CSV file

Drag and drop file here


Limit 200MB per file • CSV

Browse files

Gender

Male

Age



Hypertension

0

Heart Disease

0

Ever married

Yes


Work Type

Private

Residence Type

Rural

Average Glucose Level



Stroke Prediction App

This app predicts the probability of the onset of stroke!

Data obtained from the [Stroke Prediction Dataset](#) on Kaggle by fedesoriano.

Prediction powered by SupportVectorClassifier on sklearn.

User Input features

Awaiting CSV file to be uploaded. Currently using example input parameters (shown below).


	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_t
0	Male	30	0	0	Yes	Private	Rural

Prediction Probability

The probability of stroke onset is

0.001

See where your probability lands among the test data:





Thank you



The code for this project
can be found here:

https://github.com/JoshJingtianWang/Stroke_Prediction

